




CuTe DSL Introduce

Vicki Wang: 2025-12-06


Why CUTLASS?

Enabling innovation for SOL performance

- High-level generators leveraging compilers are popular
 - Hide and automate a lot of details
 - Get excellent performance on common use cases, but...
 - Algorithmic innovations require lower-level *abstractions*
 - Advanced features like [PDL](#) require fine grain control
- With CUTLASS
 - Available on day 0 with full control!
 - Expressive abstractions for performance in all cases
 - Modular and extensible design robust throughout GPU generations



CUTLASS in Python



- CUTLASS: A set of useful abstractions for productivity and performance at all scopes and scales
 - Open source <https://github.com/NVIDIA/cutlass>
 - Multiple entry points depending on your needs
- CuTe DSL
 - This first release makes available a *mature* low-level tensor programming model, giving access to tensor cores with full control.
 - Empower AI researchers and engineers to facilitate rapid prototyping
 - [GTC'25](#)

CUTLASS in Python

How will you get started?

pip install nvidia-cutlass-dsl



```
import cutlass
import cutlass.cute as cute

@cutte.kernel
def kernel():
    tidx, _, _ = cutlass.nvvm.thread_idx()
    if tidx == 0:
        cute.print_("Hello world")

@cutte.jit
def host():
    kernel().launch(
        grid=(1, 1, 1), block=(32, 1, 1))

host()
```



python3 hello_world.py

Writing GEMV in Python

@cute.jit / @cute.kernel


```
@cute.jit
def my_kernel(
    a_ptr: cute.Pointer,
    b_ptr: cute.Pointer,
    sfa_ptr: cute.Pointer,
    sfb_ptr: cute.Pointer,
    c_ptr: cute.Pointer,
    problem_size: tuple,
):
    # Unpack the problem size.
    m, _, k, l = problem_size

    # Create the A Tensor with (M, K, L) : (K, 1, L) layout.
    a_tensor = cute.make_tensor(
        a_ptr, cute.make_layout((m, cute.assume(k, 32), l), stride=(cute.assume(k, 32), 1, cute.assume(m * k, 32)),)
    )
    # Create the scale A Tensor with ((32, 4), REST_M), ((SF_K, 4), REST_K), (1, REST_L)) shape.
    # Details see https://docs.nvidia.com/cuda/cublas/index.html?highlight=fp4#d-block-scaling-factors-layout
    sfa_layout = blockscaled_utils.tile_atom_to_shape_SF(a_tensor.shape, sf_vec_size)
    # ... manage other parameters

    # Compute grid shape, let each block handle 128 output elements
    grid = (cute.ceil_div(c_tensor.shape[0], 128), 1, c_tensor.shape[2])

    # Launch the CUDA kernel
    kernel(a_tensor, b_tensor, sfa_tensor, sfb_tensor, c_tensor).launch(
        grid=grid,
        block=[threads_per_cta, 1, 1],
        cluster=(1, 1, 1),
    )
    return
```

- cute.jit is the entry point from host side
 - Prepare parameters for the kernel
 - Called by other DSL function



Native Implementation: Load data from global buffer to register, do
FMA and store results to global memory

Writing GEMV in Python

@cute.jit / @cute.kernel

```
@cute.kernel
def kernel(
    mA_mkl: cute.Tensor,
    mB_nkl: cute.Tensor,
    mSFA_mkl: cute.Tensor,
    mSFB_nkl: cute.Tensor,
    mC_mnl: cute.Tensor,
):
    # Get CUDA block and thread indices
    bidx, bidy, bidz = cute.arch.block_idx()
    tidx, _, _ = cute.arch.thread_idx()

    # Do other computation ...
```

- cute.kernel is the GPU kernel
 - Equal to `__global__ void kernel(...)`

Caller-Callee Compatibility


| Caller | Callee | Allowed | Compilation/Runtime |
|-----------------|-----------------|---------|--|
| Python function | @jit | ✓ | DSL runtime |
| Python function | @kernel | ✗ | N/A (error raised) |
| @jit | @jit | ✓ | Compile-time call, inlined |
| @jit | Python function | ✓ | Compile-time call, inlined, tracing-only |
| @jit | @kernel | ✓ | Dynamic call via GPU driver or runtime |
| @kernel | @jit | ✓ | Compile-time call, inlined |
| @kernel | Python function | ✓ | Compile-time call, inlined, tracing-only |
| @kernel | @kernel | ✗ | N/A (error raised) |

Writing GEMV in Python


CuTe Basic Layout

```
# Partition matrix A into local tiles (shape: [block_M, block_K, rest_M, rest_K, rest_L])
gA_mkl = cute.local_tile(
    mA_mkl, cute.slice_(mma_tiler_mnk, (None, 0, None), (None, None, None))
)
# Each block handles a tile of matrix A
# Each thread handles one row inside the tile
tAgA = gA_mkl[tidx, None, bidx, None, bidz]
```

- Do partition to split buffer before doing `cute.copy` and `cute.gemm`
- Using slice or [] to index the buffer for each block, warp, or threads



Block view of partition and slice per Batch



Thread view of partition and slice per Block

Writing GEMV in Python


Tensor and Tensor SSA

```
res = cute.zeros_like(tCgC, cutlass.Float32)
# Get the number of k tiles of the reduction loop
k_tile_cnt = gA_mkl.layout[3].shape
for k_tile in range(k_tile_cnt):
    # Slice data on k_dimension
    tAgA = gA_mkl[tidx, None, bidx, k_tile, bidz]
    tBgB = gB_nkl[0, None, bidy, k_tile, bidz]
    # Take tAgA/tBgB layout and create Rmem Tensor
    # tArA.layout: (64): (1)
    # tBrB.layout: (64): (1)
    tArA = cute.make_rmem_tensor_like(tAgA, cutlass.Float32)
    tBrB = cute.make_rmem_tensor_like(tBgB, cutlass.Float32)
    # Load NVFP4 into register (TensorSSA)
    a_val_nvfp4 = tAgA.load() # vector<64xf4E2M1FN>
    b_val_nvfp4 = tBgB.load() # vector<64xf4E2M1FN>
    # Convert float32 for computation (FFMA)
    a_val = a_val_nvfp4.to(cutlass.Float32) # vector<64xf32>
    b_val = b_val_nvfp4.to(cutlass.Float32) # vector<64xf32>
    # Store the converted values to Rmem CuTe tensors
    tArA.store(a_val)
    tBrB.store(b_val)

# Iterate over SF vector tiles and compute the scale&matmul accumulation
for i in cutlass.range_constexpr(mma_tiler_mnk[2]):
    res += tArA[i] * tBrB[i]
```

The code is simplified to show key APIs and concepts

- Tensor is the composition of an iterator and a layout
- TensorSSA thread local data modeling for CuTe Tensor in value semantics and immutable
 - Vector based with nested CuTe shape support
 - Load tensor elements as vector / store vector data into tensor
 - Operator overloading for vectorized operations



Writing GEMV in Python

Python to Python code generation: for/if/while

```
res = cute.zeros_like(tCgC, cutlass.Float32)
# Get the number of k tiles of the reduction loop
k_tile_cnt = gA_mkl.layout[3].shape
for k_tile in range(k_tile_cnt):
    # Slice data on k_dimension
    tAgA = gA_mkl[tidx, None, bidx, k_tile, bidz]
    tBgB = gB_nkl[0, None, bidy, k_tile, bidz]
    # Take tAgA/tBgB layout and create Rmem Tensor
    # tArA.layout: (64): (1)
    # tBrB.layout: (64): (1)
    tArA = cute.make_rmem_tensor_like(tAgA, cutlass.Float32)
    tBrB = cute.make_rmem_tensor_like(tBgB, cutlass.Float32)
    # Load NVFP4 into register (TensorSSA)
    a_val_nvfp4 = tAgA.load() # vector<64xf4E2M1FN>
    b_val_nvfp4 = tBgB.load() # vector<64xf4E2M1FN>
    # Convert float32 for computation (FFMA)
    a_val = a_val_nvfp4.to(cutlass.Float32) # vector<64xf32>
    b_val = b_val_nvfp4.to(cutlass.Float32) # vector<64xf32>
    # Store the converted values to Rmem CuTe tensors
    tArA.store(a_val)
    tBrB.store(b_val)

    # Iterate over SF vector tiles and compute the scale&matmul accumulation
    for i in cutlass.range_constexpr(mma_tiler_mnk[2]):
        res += tArA[i] * tBrB[i]
```

- Control flow

- CuTe DSL navigates through Python's AST and transforms each control-flow construct it encounters into a structured intermediate representation (IR). This allows users to write standard Python loops and branches, while the compiler determines whether to evaluate it at compile time if it's a native Python control flow, or to generate intermediate representation (IR) if the control flow is dynamic.
 - `cutlass.rang`: supports advanced unrolling and pipelining control
 - [Software Pipeline example](#)

Control Flow Process

| Control Flow | AST Rewrite | Python Evaluation | Scf Node Generation |
|------------------------------------|-------------|-------------------|---------------------|
| if cutlass.const_expr() | ✗ | ✓ | ✗ |
| if pred | ✓ | ✗ | ✓ |
| while cutlass.const_expr() | ✗ | ✓ | ✗ |
| while pred | ✓ | ✗ | ✓ |
| for i in cutlass.range_constexpr() | ✗ | ✓ | ✗ |
| for i in range() | ✓ | ✗ | ✓ |
| for i in cutlass.range() | ✓ | ✗ | ✓ |

Writing GEMV in Python

Software Pipeline

```
# ((atom_v, rest_v), RestK)
tBgB = tBgB[(None, mma_tile_coord_mnl[1], None, mma_tile_coord_mnl[2])]

if warp_idx == 0:
    # ///////////////////////////////////////////////////
    # MAINLOOP
    # ///////////////////////////////////////////////////
    for k_tile in cutlass.range(
        k_tile_cnt, prefetch_stages=self.num_ab_stage - 2
    ):
        # wait for AB buffer empty
        producer_handle = ab_producer.acquire_and_advance()

        # TMA load A/B
        cute.copy(...)

        cute.copy(...)

        if is_leader_cta:
            # Wait for AB buffer full
            consumer_handle = ab_consumer.wait_and_advance()

            # tCtAcc += tCrA * tCrB
            num_kblk = cute.size(tCrA, mode=[2])
            for kblk_idx in cutlass.range(num_kblk, unroll_full=True):
                kblk_crd = (None, None, kblk_idx, consumer_handle.index)

                cute.gemm(...)

            # Enable accumulate on tCtAcc after first kblock
            tiled_mma.set(tcgen05.Field.ACCUMULATE, True)

            # Async arrive AB buffer empty
            consumer_handle.release()

        # Async arrive accumulator buffer full
        if is_leader_cta:
            acc_pipeline.producer_commit(acc_producer_state)

# Pipelining TMA load A/B and MMA mainloop
#
prefetch_k_tile_cnt = cutlass.min(self.num_ab_stage - 2, k_tile_cnt)
if warp_idx == 0:
    #
    # Prefetch TMA load A/B
    #
    for k_tile_idx in cutlass.range(prefetch_k_tile_cnt, unroll=1):
        # Conditionally wait for AB buffer empty
        producer_handle = ab_producer.acquire_and_advance()

        # TMA load A/B
        cute.copy(...)

        cute.copy(...)

    peek_ab_full_status = cutlass.Boolean(False)
    if is_leader_cta:
        peek_ab_full_status = ab_consumer.try_wait()

    # Peek (try_wait) AB buffer empty for k_tile = prefetch_k_tile_cnt + k_tile + 1
    peek_ab_empty_status = ab_producer.try_acquire()

    #
    # MMA mainloop
    #
    for k_tile_idx in cutlass.range(k_tile_cnt):
        # Conditionally wait for AB buffer empty
        if k_tile_idx < k_tile_cnt - prefetch_k_tile_cnt:
            producer_handle = ab_producer.acquire_and_advance()
            peek_ab_empty_status

        # TMA load A/B
        cute.copy(...)

        cute.copy(...)

        if is_leader_cta:
            # Conditionally wait for AB buffer full
            consumer_handle = ab_consumer.wait_and_advance(peek_ab_full_status)

            # tCtAcc += tCrA * tCrB
            num_kblk = cute.size(tCrA, mode=[2])
            for kblk_idx in cutlass.range(num_kblk, unroll_full=True):
                kblk_crd = (None, None, kblk_idx, consumer_handle.index)

                cute.gemm(...)

            # Enable accumulate on tCtAcc after first kblock
            tiled_mma.set(tcgen05.Field.ACCUMULATE, True)

            # Async arrive AB buffer empty
            consumer_handle.release()

            # Peek (try_wait) AB buffer empty for k_tile = prefetch_k_tile_cnt + k_tile + 1
            peek_ab_empty_status = ab_producer.try_acquire()


            # Peek (try_wait) AB buffer full for k_tile = k_tile + 1
            peek_ab_full_status = ab_consumer.try_wait()

        # Async arrive accumulator buffer full
        if is_leader_cta:
            acc_pipeline.producer_commit(acc_producer_state)
```

Writing GEMM in Python

Use advance GPU feature: Tensor core

- Using TMA Load to load A, B, scaleA, scaleB tensors from global memory to shared memory
- Using S2T Load to load scaleA, scaleB tensors from shared memory to tensor memory
- Using tensor core instruction to do math computation
- Using T2R Load to load accumulation result from tensor memory to register
- Store results from register to global C tensor




Writing GEMM in Python

Use advance GPU feature: TMA load

```
# Setup TMA for A
tma_atom_a, tma_tensor_a = cute.nvgpu.make_tiled_tma_atom_A(
    cpasync.CopyBulkTensorTileG2S0p(tcgen05.CtaGroup.ONE),
    a_tensor,      # a_tensor.layout:(M, K, L): (K, 1, M * K)
    a_smem_layout, # a_smem_layout: S<3,4,3> o 0 o ((128,64),1,4):((256,1),0,64)
    mma_tiler_mnk, # (128, 128, 256)
    tiled_mma,     # mma_op: MmaMXF4NVF4Op(sf_dtype, (128, 128, 64),
    | | | | # CtaGroup.ONE, OperandSource.SMEM)
)

# Partition global tensor via mma_tiler shape(128, 128, 256)
# (bM, bK, RestM, RestK, RestL)
gA_mkl = cute.local_tile(
    mA_mkl, cute.slice_(self.mma_tiler, (None, 0, None)), (None, None, None)
)
# Partition global tensor via MMA op shape(128, 128, 64)
# (MMA, MMA_M, MMA_K, RestM, RestK, RestL)
tCgA = thr_mma.partition_A(gA_mkl)
# Partition global/shared tensor to fit TMA load
# ((atom_v, rest_v), STAGE)
# ((atom_v, rest_v), RestM, RestK, RestL)
tAsA, tAgA = cpasync.tma_partition(
    tma_atom_a,
    0,                      # Used for CGA config
    cute.make_layout(1),      # Used for CGA config
    cute.group_modes(sA, 0, 3), # sA.shape:((128,64),1,4,1)
    | | | | # (MMA, MMA_M, MMA_K, STAGE)
    cute.group_modes(tCgA, 0, 3), # tCgA.shape:((128,64),1,4,1,1,1)
    | | | | # (MMA, MMA_M, MMA_K, RestM, RestK, RestL)
)
# Execute k_tile loop
for k_tile in range(k_tile_cnt):
    # TMA Load A Tensor
    cute.copy(
        tma_atom_a,
        tAgA[(None, k_tile)],
        tAsA[(None, ab_empty.index)],
        tma_bar_ptr=ab_empty.barrier, # TMA barrier pointer
    )
```

Host codes to setup
TMA atom and TMA
coord tensor
(in @cute.jit)



The code is simplified to show key APIs and concepts

Writing GEMM in Python

Use advance GPU feature: TMA load


```
# Setup TMA for A
tma_atom_a, tma_tensor_a = cute.nvgpu.make_tiled_tma_atom_A(
    cpasync.CopyBulkTensorTileG2S0p(tcgen05.CtaGroup.ONE),
    a_tensor,      # a_tensor.layout:(M, K, L): (K, 1, M * K)
    a_smem_layout, # a_smem_layout: S<3,4,3> o 0 o ((128,64),1,4):((256,1),0,64)
    mma_tiler_mnk, # (128, 128, 256)
    tiled_mma,     # mma_op: MmaMXF4NVF4Op(sf_dtype, (128, 128, 64),
    | | | | # CtaGroup.ONE, OperandSource.SMEM)
)

# Partition global tensor via mma_tiler shape(128, 128, 256)
# (bM, bK, RestM, RestK, RestL)
gA_mkl = cute.local_tile(
    mA_mkl, cute.slice_(self.mma_tiler, (None, 0, None)), (None, None, None)
)
# Partition global tensor via MMA op shape(128, 128, 64)
# (MMA, MMA_M, MMA_K, RestM, RestK, RestL)
tCgA = thr_mma.partition_A(gA_mkl)
# Partition global/shared tensor to fit TMA load
# ((atom_v, rest_v), STAGE)
# ((atom_v, rest_v), RestM, RestK, RestL)
tAsA, tAgA = cpasync.tma_partition(
    tma_atom_a,
    0,                      # Used for CGA config
    cute.make_layout(1),      # Used for CGA config
    cute.group_modes(sA, 0, 3), # sA.shape:((128,64),1,4,1)
    | | | | # (MMA, MMA_M, MMA_K, STAGE)
    cute.group_modes(tCgA, 0, 3), # tCgA.shape:((128,64),1,4,1,1,1)
    | | | | # (MMA, MMA_M, MMA_K, RestM, RestK, RestL)
)

# Execute k_tile loop
for k_tile in range(k_tile_cnt):
    # TMA Load A Tensor
    cute.copy(
        tma_atom_a,
        tAgA[(None, k_tile)],
        tAsA[(None, ab_empty.index)],
        tma_bar_ptr=ab_empty.barrier, # TMA barrier pointer
    )
```

Host codes to setup
TMA atom and TMA
coord tensor
(in @cute.jit)

GPU codes: Partition
Gmem and Smem buffer
to prepare inputs for TMA
load instruction




Writing GEMM in Python

Use advance GPU feature: TMA load

Host codes to setup
TMA atom and TMA
coord tensor
(in @cute.jit)

GPU codes: Partition
Gmem and Smem buffer
to prepare inputs for TMA
load instruction

GPU codes: Call TMA
load instruction
For each k tile




Writing GEMM in Python

Use advance GPU feature: S2T copy

```
#  
# Partition for S2T copy of SFA/SFB  
#  
# Make S2T CopyAtom  
copy_atom_s2t = cute.make_copy_atom(  
    tcgen05.Cp4x32x128b0p(tcgen05.CtaGroup.ONE),  
    sf_dtype,  
)  
tiled_copy_s2t_sfa = tcgen05.make_s2t_copy(copy_atom_s2t, tCtSFA)  
thr_copy_s2t_sfa = tiled_copy_s2t_sfa.get_slice(0)  
  
# ((ATOM_V, REST_V), Rest_Tiler, MMA_MN, MMA_K, STAGE)  
tCsSFA_s2t = thr_copy_s2t_sfa.partition_S(ssFA)  
# ((ATOM_V, REST_V), Rest_Tiler, MMA_MN, MMA_K)  
tCtSFA_s2t = thr_copy_s2t_sfa.partition_D(tCtSFA)  
  
for k_tile in cutlass.range(k_tile_cnt):  
    # ... copy data from global to shared memory  
  
    # Copy SFA/SFB to tmem  
    s2t_stage_coord = (None, None, None, None, ab_full.index)  
    tCsSFA_s2t_staged = tCsSFA_s2t[s2t_stage_coord]  
    cute.copy(  
        tiled_copy_s2t_sfa,  
        tCsSFA_s2t_staged,  
        tCtSFA_s2t,  
)
```

Construct the Tiled S2T
copy with proper
configure



The code is simplified to show key APIs and concepts

Writing GEMM in Python


Use advance GPU feature: S2T copy

```
#  
# Partition for S2T copy of SFA/SFB  
  
#  
# Make S2T CopyAtom  
copy_atom_s2t = cute.make_copy_atom(  
    tcgen05.Cp4x32x128b0p(tcgen05.CtaGroup.ONE),  
    sf_dtype,  
)  
tiled_copy_s2t_sfa = tcgen05.make_s2t_copy(copy_atom_s2t, tCtSFA)  
thr_copy_s2t_sfa = tiled_copy_s2t_sfa.get_slice(0)  
  
# ((ATOM_V, REST_V), Rest_Tiler, MMA_MN, MMA_K, STAGE)  
tCsSFA_s2t = thr_copy_s2t_sfa.partition_S(sSFA)  
# ((ATOM_V, REST_V), Rest_Tiler, MMA_MN, MMA_K)  
tCtSFA_s2t = thr_copy_s2t_sfa.partition_D(tCtSFA)  
  
for k_tile in cutlass.range(k_tile_cnt):  
    # ... copy data from global to shared memory  
  
    # Copy SFA/SFB to tmem  
    s2t_stage_coord = (None, None, None, None, ab_full.index)  
    tCsSFA_s2t_staged = tCsSFA_s2t[s2t_stage_coord]  
    cute.copy(  
        tiled_copy_s2t_sfa,  
        tCsSFA_s2t_staged,  
        tCtSFA_s2t,  
)
```

Construct the Tiled S2T copy with proper configure

Partition the shared memory and tensor memory for S2T copy

The code is simplified to show key APIs and concepts



Writing GEMM in Python

Use advance GPU feature: S2T copy


```
#  
# Partition for S2T copy of SFA/SFB  
  
#  
# Make S2T CopyAtom  
copy_atom_s2t = cute.make_copy_atom(  
    tcgen05.Cp4x32x128b0p(tcgen05.CtaGroup.ONE),  
    sf_dtype,  
)  
tiled_copy_s2t_sfa = tcgen05.make_s2t_copy(copy_atom_s2t, tCtSFA)  
thr_copy_s2t_sfa = tiled_copy_s2t_sfa.get_slice(0)  
  
# ((ATOM_V, REST_V), Rest_Tiler, MMA_MN, MMA_K, STAGE)  
tCsSFA_s2t = thr_copy_s2t_sfa.partition_S(sSFA)  
# ((ATOM_V, REST_V), Rest_Tiler, MMA_MN, MMA_K)  
tCtSFA_s2t = thr_copy_s2t_sfa.partition_D(tCtSFA)  
  
for k_tile in cutlass.range(k_tile_cnt):  
    # ... copy data from global to shared memory  
  
    # Copy SFA/SFB to tmem  
    s2t_stage_coord = (None, None, None, None, ab_full.index)  
    tCsSFA_s2t_staged = tCsSFA_s2t[s2t_stage_coord]  
    cute.copy(  
        tiled_copy_s2t_sfa,  
        tCsSFA_s2t_staged,  
        tCtSFA_s2t,  
)
```

The code is simplified to show key APIs and concepts

Construct the Tiled S2T copy with proper configure

Partition the shared memory and tensor memory for S2T copy

Call S2T copy for each kBlock



Writing GEMM in Python


Use advance GPU feature: Tensor core

```
# Pick the right tensor core operation
mma_op = tcgen05.MmaMXF4NVF4Op(
    sf_dtype,
    (128, 128, 64), # (MMA, MMA_M, MMA_K)
    tcgen05.CtaGroup.ONE,
    tcgen05.OperandSource.SMEM,
)
tiled_mma = cute.make_tiled_mma(mma_op)

# (MMA, MMA_M, MMA_K, STAGE)
tCrA = tiled_mma.make_fragment_A(sA)
# (MMA, MMA_N, MMA_K, STAGE)
tCrB = tiled_mma.make_fragment_B(sB)
# (MMA, MMA_M, MMA_N)
tCtAcc = tiled_mma.make_fragment_C(mma_tiler_mnk[:2])

# For each k block, set the scale factor tensors iterator in tiled_mma
# and perform the blockscaled GEMM tensor-core operation.
num_kblocks = cute.size(tCrA, mode=[2])
for kblock_idx in cutlass.range(num_kblocks, unroll_full=True):
    tiled_mma.set(
        tcgen05.Field.SFA,
        gCtSFA[kblock_idx].iterator,
    )
    tiled_mma.set(
        tcgen05.Field.SFB,
        gCtSFB[kblock_idx].iterator,
    )
    cute.gemm(
        tiled_mma,
        tCtAcc,
        tCrA[kblock_idx],
        tCrB[kblock_idx],
        tCtAcc,
    )
```

Construct the Tiled
MMA with right
TensorCore configure



The code is simplified to show key APIs and concepts

Writing GEMM in Python

Use advance GPU feature: Tensor core


```
# Pick the right tensor core operation
mma_op = tcgen05.MmaMXF4NVF40p(
    sf_dtype,
    (128, 128, 64), # (MMA, MMA_M, MMA_K)
    tcgen05.CtaGroup.ONE,
    tcgen05.OperandSource.SMEM,
)
tiled_mma = cute.make_tiled_mma(mma_op)

# (MMA, MMA_M, MMA_K, STAGE)
tCrA = tiled_mma.make_fragment_A(sA)
# (MMA, MMA_N, MMA_K, STAGE)
tCrB = tiled_mma.make_fragment_B(sB)
# (MMA, MMA_M, MMA_N)
tCtAcc = tiled_mma.make_fragment_C(mma_tiler_mnk[:2])

# For each k block, set the scale factor tensors iterator in tiled_mma
# and perform the blockscaled GEMM tensor-core operation.
num_kblocks = cute.size(tCrA, mode=[2])
for kblock_idx in cutlass.range(num_kblocks, unroll_full=True):
    tiled_mma.set(
        tcgen05.Field.SFA,
        gCtSFA[kblock_idx].iterator,
    )
    tiled_mma.set(
        tcgen05.Field.SFB,
        gCtSFB[kblock_idx].iterator,
    )
    cute.gemm(
        tiled_mma,
        tCtAcc,
        tCrA[kblock_idx],
        tCrB[kblock_idx],
        tCtAcc,
    )
```

Construct the Tiled
MMA with right
TensorCore configure

Construct A, B, C operands
for the TensorCore, A and B
come from Smem and Acc
comes from TMEM



The code is simplified to show key APIs and concepts

Writing GEMM in Python

Use advance GPU feature: Tensor core

```
# Pick the right tensor core operation
mma_op = tcgen05.MmaMXF4NVF40p(
    sf_dtype,
    (128, 128, 64), # (MMA, MMA_M, MMA_K)
    tcgen05.CtaGroup.ONE,
    tcgen05.OperandSource.SMEM,
)
tiled_mma = cute.make_tiled_mma(mma_op)


# (MMA, MMA_M, MMA_K, STAGE)
tCrA = tiled_mma.make_fragment_A(sA)
# (MMA, MMA_N, MMA_K, STAGE)
tCrB = tiled_mma.make_fragment_B(sB)
# (MMA, MMA_M, MMA_N)
tCtAcc = tiled_mma.make_fragment_C(mma_tiler_mnk[:2])

# For each k block, set the scale factor tensors iterator in tiled_mma
# and perform the blockscaled GEMM tensor-core operation.
num_kblocks = cute.size(tCrA, mode=[2])
for kblock_idx in cutlass.range(num_kblocks, unroll_full=True):
    tiled_mma.set(
        tcgen05.Field.SFA,
        gCtSFA[kblock_idx].iterator,
    )
    tiled_mma.set(
        tcgen05.Field.SFB,
        gCtSFB[kblock_idx].iterator,
    )
    cute.gemm(
        tiled_mma,
        tCtAcc,
        tCrA[kblock_idx],
        tCrB[kblock_idx],
        tCtAcc,
```

Construct the Tiled
MMA with right
TensorCore configure

Construct A, B, C operands
for the TensorCore, A and B
come from Smem and Acc
comes from TMEM

For each TensorCore in the
kBlock, update scaling factor
iterators and do the
accumulation



The code is simplified to show key APIs and concepts


Writing GEMM in Python

Use advance GPU feature: T2R copy

```
#  
# Epilogue  
# Partition for epilogue  
#  
op = tcgen05.Ld32x32b0p(tcgen05.Repetition.x128, tcgen05.Pack.NONE)  
copy_atom_t2r = cute.make_copy_atom(op, cutlass.Float32)  
tiled_copy_t2r = tcgen05.make_tmem_copy(copy_atom_t2r, tCtAcc)  
thr_copy_t2r = tiled_copy_t2r.get_slice(tidx)  
  
# (T2R_M, T2R_N, EPI_M, EPI_N)  
tTR_tAcc = thr_copy_t2r.partition_S(tCtAcc)  
# (T2R_M, T2R_N, EPI_M, EPI_N, RestM, RestN, RestL)  
tTR_gC = thr_copy_t2r.partition_D(tCgC)  
  
# (T2R_M, T2R_N, EPI_M, EPI_N)  
tTR_rAcc = cute.make_rmem_tensor(  
    tTR_gC[None, None, None, None, 0, 0, 0].shape, cutlass.Float32  
)  
# Copy accumulator to register  
cute.copy(tiled_copy_t2r, tTR_tAcc, tTR_rAcc)
```

Construct the Tiled T2R
copy with proper
configure

The code is simplified to show key APIs and concepts




Writing GEMM in Python

Use advance GPU feature: T2R copy

```
#  
# Epilogue  
# Partition for epilogue  
#  
op = tcgen05.Ld32x32b0p(tcgen05.Repetition.x128, tcgen05.Pack.NONE) Construct the Tiled T2R  
copy_atom_t2r = cute.make_copy_atom(op, cutlass.Float32) copy with proper  
tiled_copy_t2r = tcgen05.make_tmem_copy(copy_atom_t2r, tCtAcc) configure  
thr_copy_t2r = tiled_copy_t2r.get_slice(tidx)  
  
# (T2R_M, T2R_N, EPI_M, EPI_M)  
tTR_tAcc = thr_copy_t2r.partition_S(tCtAcc)  
# (T2R_M, T2R_N, EPI_M, EPI_N, RestM, RestN, RestL)  
tTR_gC = thr_copy_t2r.partition_D(tCgC)  
  
# (T2R_M, T2R_N, EPI_M, EPI_N)  
tTR_rAcc = cute.make_rmem_tensor(  
    tTR_gC[None, None, None, None, 0, 0, 0].shape, cutlass.Float32  
)  
  
# Copy accumulator to register  
cute.copy(tiled_copy_t2r, tTR_tAcc, tTR_rAcc)
```

Slice via thread view as T2R
is SIMD instruction



The code is simplified to show key APIs and concepts

Writing GEMM in Python

Use advance GPU feature: T2R copy


```
#  
# Epilogue  
# Partition for epilogue  
#  
op = tcgen05.Ld32x32b0p(tcgen05.Repetition.x128, tcgen05.Pack.NONE) Construct the Tiled T2R  
copy_atom_t2r = cute.make_copy_atom(op, cutlass.Float32) copy with proper  
tiled_copy_t2r = tcgen05.make_tmem_copy(copy_atom_t2r, tCtAcc) configure  
thr_copy_t2r = tiled_copy_t2r.get_slice(tidx)  
  
# (T2R_M, T2R_N, EPI_M, EPI_M)  
tTR_tAcc = thr_copy_t2r.partition_S(tCtAcc) Slice via thread view as T2R  
# (T2R_M, T2R_N, EPI_M, EPI_N, RestM, RestN, RestL) is SIMT instruction  
tTR_gC = thr_copy_t2r.partition_D(tCgC)  
  
# (T2R_M, T2R_N, EPI_M, EPI_N)  
tTR_rAcc = cute.make_rmem_tensor(  
    tTR_gC[None, None, None, None, 0, 0, 0].shape, cutlass.Float32)  
)  
  
# Copy accumulator to register  
cute.copy(tiled_copy_t2r, tTR_tAcc, tTR_rAcc)
```

The code is simplified to show key APIs and concepts

Construct the Tiled T2R
copy with proper
configure

Slice via thread view as T2R
is SIMT instruction

Partition the tensor
memory and global
tensor and make register
tensor based on global
tensor shape for T2R
copy



Writing GEMM in Python

Use advance GPU feature: T2R copy


```
#  
# Epilogue  
# Partition for epilogue  
#  
op = tcgen05.Ld32x32b0p(tcgen05.Repetition.x128, tcgen05.Pack.NONE) Construct the Tiled T2R  
copy_atom_t2r = cute.make_copy_atom(op, cutlass.Float32) copy with proper  
tiled_copy_t2r = tcgen05.make_tmem_copy(copy_atom_t2r, tCtAcc) configure  
thr_copy_t2r = tiled_copy_t2r.get_slice(tidx)  
  
# (T2R_M, T2R_N, EPI_M, EPI_N)  
tTR_tAcc = thr_copy_t2r.partition_S(tCtAcc)  
# (T2R_M, T2R_N, EPI_M, EPI_N, RestM, RestN, RestL)  
tTR_gC = thr_copy_t2r.partition_D(tCgC)  
  
# (T2R_M, T2R_N, EPI_M, EPI_N)  
tTR_rAcc = cute.make_rmem_tensor(  
    tTR_gC[None, None, None, None, 0, 0, 0].shape, cutlass.Float32  
)  
  
# Copy accumulator to register  
cute.copy(tiled_copy_t2r, tTR_tAcc, tTR_rAcc)
```

The code is simplified to show key APIs and concepts

Construct the Tiled T2R
copy with proper
configure
Slice via thread view as T2R
is SIMD instruction

Partition the tensor
memory and global
tensor and make register
tensor based on global
tensor shape for T2R
copy

Call the T2R copy



Writing GEMM in Python

Cute Programming Model

- Typical code sequence (APIs) used to do copy and gemm

- Copy
 - local_tile (tiling global tensor)
 - partition_S, partition_D Or tma_partition
 - cute.copy
- Mma
 - make_fragment_A
 - make_fragment_B
 - make_fragment_C
 - cute.gemm

- Variable Naming Convention

- tXaY : Partitioning pattern tX applied to tensor aY
 - *a*: is r for register, g for Gmem, s for Smem, c for coordinate, p for predicate, t for Tmem
- For example:

```
# Define thread layouts (static)
tC = make_layout((16, 16), stride=(16, 1));
# Partition gC (M,N) by the tile of tC
tCgC = cute.local_partition(gC, tC, tidx, (1, 1));
```

Writing GEMM in Python

@cute.struct - Customized C struct like data types

- cute.struct decorator
 - Transform a Python class into a memory-mapped structure with precise control over memory layout, alignment and offsets
 - Support scalar, MemRange or nested struct as data members
 - Allow customized data alignment which is essential for better performance
- SmemAllocator
 - Manage shared memory allocation of raw bytes, numeric types, and tensors with static layout

```
@cute.struct
class SharedStorage:
    ab_mbar_ptr: cute.struct.MemRange[cutlass.Int64, self.num_ab_stage * 2]
    acc_mbar_ptr: cute.struct.MemRange[cutlass.Int64, self.num_acc_stage * 2]
    tmem_holding_buf: cutlass.Int32

smem = utils.SmemAllocator()
storage = smem.allocate(SharedStorage)
# (MMA, MMA_M, MMA_K, STAGE)
sA = smem.allocate_tensor(
    element_type=ab_dtype,
    layout=a_smem_layout_staged.outer,
    byte_alignment=128,
    swizzle=a_smem_layout_staged.inner,
)
```

| offset | SharedStorage | alignment |
|-------------------------|------------------|-----------|
| 0 | ab_mbar_ptr | 8 |
| + 8 * 2 * num_ab_stage | acc_mbar_ptr | 8 |
| + 8 * 2 * num_acc_stage | tmem_holding_buf | 4 |

Natural alignment per dtype

Writing GEMM in Python

Pipeline – warp specialization

```

#
# Initialize mainloop ab_pipeline, acc_pipeline and their states
#
ab_pipeline_producer_group = pipeline.CooperativeGroup(pipeline.Agent.Thread)
ab_pipeline_consumer_group = pipeline.CooperativeGroup(pipeline.Agent.Thread, 1)
ab_producer, ab_consumer = pipeline.PipelineTmaUmma.create(
    barrier_storage=storage.ab_mbar_ptr.data_ptr(),
    num_stages=self.num_ab_stage,
    producer_group=ab_pipeline_producer_group,
    consumer_group=ab_pipeline_consumer_group,
    tx_count=self.num_tma_load_bytes,
).make_participants()
acc_producer, acc_consumer = pipeline.PipelineUmmaAsync.create(
    barrier_storage=storage.acc_mbar_ptr.data_ptr(),
    num_stages=self.num_acc_stage,
    producer_group=ab_pipeline_producer_group,
    consumer_group=pipeline.CooperativeGroup(
        pipeline.Agent.Thread,
        self.threads_per_cta,
    ),
).make_participants()

```

Async pipeline allow producer and consumer work concurrently for best performance

Between TMA and Tensor Core (PipelineTmaUmma)

```

if warp_idx == 0:
    # Execute k_tile loop
    for k_tile in range(k_tile_cnt):
        # Wait for AB buffer empty
        ab_empty = ab_producer.acquire_and_advance()

        # TMA Load A/B/SFA/SFB
        cute.copy(...
        )
        cute.copy(...
        )
        cute.copy(...
        )
        cute.copy(...
        )

# Wait for AB buffer full
ab_full = ab_consumer.wait_and_advance()
# tCtAcc += tCrA * tCrSFA * tCrB * tCrSFB
num_kblocks = cute.size(tCrA, mode=[2])
for kblock_idx in cutlass.range(num_kblocks, unroll_full=True):
    kblock_coord = (...)

    # Set SFA/SFB tensor to tiled_mma
    sf_kblock_coord = (None, None, kblock_idx)
    tiled_mma.set(...
    )
    tiled_mma.set(...
    )

    cute.gemm(...
    )
    # Enable accumulate on tCtAcc after first kblock
    tiled_mma.set(tcgen05.Field.ACUMULATE, True)
# Async arrive AB buffer empty
ab_full.release()

```

Between Tensor Core and Epilogue (PipelineUmmaAsync)

```

# Wait for accumulator buffer empty
acc_empty = acc_producer.acquire_and_advance()
# Set ACCUMULATE field to False for the first k_tile iteration
tiled_mma.set(tcgen05.Field.ACUMULATE, False)
# Execute k_tile loop
for k_tile in range(k_tile_cnt):
    # Wait for AB buffer empty
    ab_empty = ab_producer.acquire_and_advance()

    # TMA Load A/B/SFA/SFB
    cute.copy(...
    )
    # ...
    # Wait for AB buffer full
    ab_full = ab_consumer.wait_and_advance()

    # tCtAcc += tCrA * tCrSFA * tCrB * tCrSFB
    num_kblocks = cute.size(tCrA, mode=[2])
    for kblock_idx in cutlass.range(num_kblocks, unroll_full=True):
        # ...
        cute.gemm(...
        )
        # Enable accumulate on tCtAcc after first kblock
        tiled_mma.set(tcgen05.Field.ACUMULATE, True)

    # Async arrive AB buffer empty
    ab_full.release()
    acc_empty.commit()

# Wait for accumulator buffer full
acc_full = acc_consumer.wait_and_advance()

# Copy accumulator to register
cute.copy(tiled_copy_t2r, tTR_tAcc, tTR_rAcc)
acc_vec = epilogue_op(tTR_rAcc.load().to(c_dtype))
tTR_rC.store(acc_vec)
# Store C to global memory
cute.copy(simt_atom, tTR_rC, tTR_gC)

acc_full.release()

```

Writing Dual GEMM in Python

Static argument vs. Dynamic argument ([doc](#))


- **Static arguments** hold values that are known at compile time. It is not included in the generated JIT function signature.
- **Dynamic arguments** hold values that are only known at runtime (default behavior)
- `cutlass.Constexpr` can be used to specify a static argument.

```
# GPU device kernel
@cute.kernel
def kernel(
    self,
    tiled_mma: cute.TiledMma,
    tma_atom_a: cute.CopyAtom,
    # ...
    epilogue_op: cutlass.Constexpr = lambda x: x
    * (1.0 / (1.0 + cute.math.exp(-x, fastmath=True))),
):
    # (T2R_M, T2R_N, EPI_M, EPI_N)
    tTR_rAcc1 = cute.make_rmem_tensor(
        tTR_gC[None, None, None, None, 0, 0, 0].shape, cutlass.Float32
    )
    # (T2R_M, T2R_N, EPI_M, EPI_N)
    tTR_rAcc2 = cute.make_rmem_tensor(
        tTR_gC[None, None, None, None, 0, 0, 0].shape, cutlass.Float32
    )

    # Copy accumulator from Tmem to Rmem
    cute.copy(tiled_copy_t2r, tTR_tAcc1, tTR_rAcc1)
    cute.copy(tiled_copy_t2r, tTR_tAcc2, tTR_rAcc2)

    # Silu activation on acc1 and multiply with acc2
    acc_vec1 = epilogue_op(tTR_rAcc1.load())
    acc_vec2 = tTR_rAcc2.load()
    acc_vec = acc_vec1 * acc_vec2
```

Constexpr to define swiglu
fusion in epilogue



Writing Dual GEMM in Python

[Tensor and TensorSSA \(doc\)](#)


- [TensorSSA](#) thread local data modeling for CuTe Tensor in value semantics and immutable
 - Vector based with nested CuTe shape support
 - Load tensor elements as vector / store vector data into tensor
 - Operator overloading for vectorized operations

```
# GPU device kernel
@cute.kernel
def kernel(
    self,
    tiled_mma: cute.TiledMma,
    tma_atom_a: cute.CopyAtom,
    # ...
    epilogue_op: cutlass.Constexpr = lambda x: x
        * (1.0 / (1.0 + cute.math.exp(-x, fastmath=True))),
):
    # (T2R_M, T2R_N, EPI_M, EPI_N)
    tTR_rAcc1 = cute.make_rmem_tensor(
        tTR_gC[None, None, None, None, 0, 0, 0].shape, cutlass.Float32
    )
    # (T2R_M, T2R_N, EPI_M, EPI_N)
    tTR_rAcc2 = cute.make_rmem_tensor(
        tTR_gC[None, None, None, None, 0, 0, 0].shape, cutlass.Float32
    )

    # Copy accumulator from Tmem to Rmem
    cute.copy(tiled_copy_t2r, tTR_tAcc1, tTR_rAcc1)
    cute.copy(tiled_copy_t2r, tTR_tAcc2, tTR_rAcc2)

    # Silu activation on acc1 and multiply with acc2
    acc_vec1 = epilogue_op(tTR_rAcc1.load())
    acc_vec2 = tTR_rAcc2.load()
    acc_vec = acc_vec1 * acc_vec2
```

Constexpr to define swiglu fusion in epilogue




Load data to register, do swiglu on Acc1 and then do elementwise multiply for Acc1 and Acc2

Writing Group GEMM in Python

Update TMA on the fly

- CTA will figure which group to compute on device and thus need to update TMA desc from device before doing `cute.copy`
 - **TensorMapManager**: provides a set of utils to manage TensorMap operations, including initialization and updates
 - **Step1**
 - Each CTA gets its own Gmem tensormap pointers for A/B/SFA/SFB tensors

```
tensormap_manager = utils.TensorMapManager(  
    utils.TensorMapUpdateMode.SMEM,  
    128,  
)  
tensormap_a_gmem_ptr = tensormap_manager.get_tensormap_ptr(  
    tensormaps[(bidz, 0, None)].iterator  
)  
tensormap_b_gmem_ptr = tensormap_manager.get_tensormap_ptr(  
    tensormaps[(bidz, 1, None)].iterator  
)  
tensormap_sfa_gmem_ptr = tensormap_manager.get_tensormap_ptr(  
    tensormaps[(bidz, 2, None)].iterator  
)  
tensormap_sfb_gmem_ptr = tensormap_manager.get_tensormap_ptr(  
    tensormaps[(bidz, 3, None)].iterator  
)
```




Writing Group GEMM in Python

Update TMA on the fly

- CTA will figure which group to compute on device and thus need to update TMA desc from device before doing `cute.copy`
 - **TensorMapManager**: provides a set of utils to manage TensorMap operations, including initialization and updates
 - **Step2**
 - Construct tma desc in shared memory
 - Copying tensormap (with fake globalAddress, globalDim, globalStrides) from const memory to Smem

```
if warp_idx == 0:  
    tensormap_manager.init_tensormap_from_atom(  
        tma_atom_a, tensormap_a_smem_ptr, 0  
    )  
    tensormap_manager.init_tensormap_from_atom(  
        tma_atom_b, tensormap_b_smem_ptr, 0  
    )  
    tensormap_manager.init_tensormap_from_atom(  
        tma_atom_sfa, tensormap_sfa_smem_ptr, 0  
    )  
    tensormap_manager.init_tensormap_from_atom(  
        tma_atom_sfb, tensormap_sfb_smem_ptr, 0  
    )
```




Writing Group GEMM in Python

Update TMA on the fly

- CTA will figure which group to compute on device and thus need to update TMA desc from device before doing `cute.copy`
 - **TensorMapManager**: provides a set of utils to manage TensorMap operations, including initialization and updates
 - **Step3**
 - Update tensormaps' globalAddress, globalDim, globalStrides with real global memory addresses and problem sizes.
 - Store tma desc to global tensormap space

```
tensormap_manager.update_tensormap(  
    (  
        real_tensor_a,  
        real_tensor_b,  
        real_tensor_sfa,  
        real_tensor_sfb,  
    ),  
    (tma_atom_a, tma_atom_b, tma_atom_sfa, tma_atom_sfb),  
    (  
        tensormap_a_gmem_ptr,  
        tensormap_b_gmem_ptr,  
        tensormap_sfa_gmem_ptr,  
        tensormap_sfb_gmem_ptr,  
    ),  
    0, # tma warp id  
    (  
        tensormap_a_smem_ptr,  
        tensormap_b_smem_ptr,  
        tensormap_sfa_smem_ptr,  
        tensormap_sfb_smem_ptr,  
    ),  
)
```




Writing Group GEMM in Python

Update TMA on the fly

- CTA will figure which group to compute on device and thus need to update TMA desc from device before doing `cute.copy`
 - **TensorMapManager**: provides a set of utils to manage TensorMap operations, including initialization and updates
 - **Step4**
 - Ensure TMA desc are stored in Gmem

```
# Let warp 0 initialize tensormap
if warp_idx == 0:
    tensormap_manager.fence_tensormap_update(tensormap_a_gmem_ptr)
    tensormap_manager.fence_tensormap_update(tensormap_b_gmem_ptr)
    tensormap_manager.fence_tensormap_update(tensormap_sfa_gmem_ptr)
    tensormap_manager.fence_tensormap_update(tensormap_sfb_gmem_ptr)
```



Compile and Run

Cache

```
# Global cache for compiled kernels (keyed by group size)
_compiled_kernel_cache = {}

# This function is used to compile the kernel once and cache it and then allow users to
# run the kernel multiple times to get more accurate timing results.

def compile_kernel(problem_sizes):
    global _compiled_kernel_cache

    # Convert problem_sizes list to a hashable tuple for use as dictionary key
    cache_key = f"{len(problem_sizes)}"

    # Check if we already have a compiled kernel for these problem sizes
    if cache_key in _compiled_kernel_cache:
        return _compiled_kernel_cache[cache_key]

    cute_ptr_of_tensor_of_problem_sizes = make_ptr(
        cutlass.Int32, 0, cute.AddressSpace.gmem, assumed_align=16,
    )Use nullptr to compile
    # Create fake tensors and pointers to compile the kernel
    compiled_func = cute.compile(
        my_kernel,
        cute_ptr_of_tensor_of_problem_sizes,
        cute_ptr_of_tensor_of_abc_ptrs,
        cute_ptr_of_tensor_of_sfafb_ptrs,
        cute_ptr_of_tensor_of_tensormap,
        total_num_clusters,
        problem_sizes,
        num_groups
    )
    # Store compiled kernel in cache with problem_sizes as key
    _compiled_kernel_cache[cache_key] = compiled_func
    return compiled_func
```

Cache the kernel
based on group
sizes

Check if the cache
contain the cubin

```
def custom_kernel(data: input_t) -> output_t:
    abc_tensors, _, sfafb_reordered_tensors, problem_sizes = data

    compiled_func = compile_kernel(problem_sizes)

    # Extract raw data pointers from all input tensors for each group
    # These will be passed to the GPU kernel to access the actual tensor data
    abc_ptrs = []
    sfafb_ptrs = []
    for i, ((a, b, c), (sfa_reordered, sfb_reordered), (m, n, k, l)) in
        enumerate(zip(abc_tensors, sfafb_reordered_tensors, problem_sizes)):
        # Store pointers to A, B, and C matrices for this group
        abc_ptrs.append((a.data_ptr(), b.data_ptr(), c.data_ptr()))
        # Store pointers to scale factor tensors for this group
        sfafb_ptrs.append((sfa_reordered.data_ptr(), sfb_reordered.data_ptr()))

    # ... codes to prepare arguments of the kernel

    # Create CuTe pointers to the metadata tensors that will be passed to the kernel
    # These allow the GPU kernel to read problem sizes and tensor pointers
    cute_ptr_of_tensor_of_abc_ptrs = make_ptr(
        cutlass.Int64,
        tensor_of_abc_ptrs.data_ptr(),Pass the real data when run
        cute.AddressSpace.gmem,the kernel
        assumed_align=16,
    )
    # ... codes to prepare arguments of the kernel

    # Launch the JIT-compiled GPU kernel with all prepared data
    # The kernel will perform block-scaled group GEMM: C = A * SFA * B * SFB for all groups
    compiled_func(
        cute_ptr_of_tensor_of_problem_sizes, # Pointer to problem sizes array
        cute_ptr_of_tensor_of_abc_ptrs, # Pointer to ABC tensor pointers array
        cute_ptr_of_tensor_of_sfafb_ptrs, # Pointer to scale factor pointers array
        cute_ptr_of_tensor_of_tensormap, # Pointer to tensormap buffer
        total_num_clusters, # Total number of CTAs to launch
        problem_sizes, # Problem sizes list (for host-side processing)
        num_groups,# Number of groups in this batch
    )
```

Debug in Python

cute.print/cute.print tensor

- **print:** Can only show static values at compile time
- **cute.printf:** Can display both static and dynamic values at runtime
- **cute.print_tensor:** Dump contents of the tensor
 - Only support Float16/Float32/Float64 data type

```
# (bM, bN, RestM, RestN, RestL)
gC_mnl = cute.local_tile(
    mC_mnl, cute.slice_(self.mma_tiler, (None, None, 0)), (None, None, None)
)
print(f"mC_mnl: {mC_mnl}")
print(f"gC_mnl: {gC_mnl}")
if tidx == 0:
    cute.printf("mC_mnl: {}", mC_mnl)
    cute.printf("gC_mnl: {}", gC_mnl)
    cute.print_tensor(mC_mnl)
    cute.print_tensor(gC_mnl[None, 0, 0, 0, 0])
    cute.print_tensor(gC_mnl[None, 0, 0, 0, 0], verbose=True)
```

```
mC_mnl: tensor<ptr<f16, gmem, align<16>> o ({div=32},?,?):(?,1,?)>
gC_mnl: tensor<ptr<f16, gmem, align<16>> o (128,128,?,?):(?,1,{div=128},128,?)>
mC_mnl: raw_ptr(0x0000738313008000: f16, gmem, align<16>) o (128,128,1):(128,1,16384)
gC_mnl: raw_ptr(0x0000738313008000: f16, gmem, align<16>) o (128,128,1,1,1):(128,1,16384,128,16384)
tensor(raw_ptr(0x0000738313008000: f16, gmem, align<16>) o (128,128,1):(128,1,16384), data=
[[[-0.017166, -0.003027,  0.513672, ..., -1.078125,  0.255615,  0.859375, ],
 [ 0.047668, -0.700684,  0.719238, ..., -1.219727,  0.462402,  1.126953, ],
 [-1.769531,  0.977051,  0.722656, ..., -1.695312,  0.187012, -0.850586, ],
 ...
 [ 0.396484,  0.466309, -0.545898, ..., -1.271484,  0.546875, -0.288818, ],
 [-0.044281, -1.361328, -0.631836, ..., -0.820801,  1.089844,  0.040985, ],
 [-0.324707, -1.977539,  1.054688, ...,  0.364746,  0.242188, -1.424805, ]])
tensor(raw_ptr(0x0000738313008000: f16, gmem, align<16>) o (128):(128), data=
[-0.017166, ],
[ 0.047668, ],
[-1.769531, ],
...
[ 0.396484, ],
[-0.044281, ],
[-0.324707, ])
tensor(raw_ptr(0x0000738313008000: f16, gmem, align<16>) o (128):(128), data=
(0)=-0.017166
(1)= 0.047668
(2)=-1.769531
(3)= 0.019379
(4)=-0.750488
(5)=-0.035400
(6)=-1.210938
(7)= 1.065430
(8)=-0.353271
(9)= 0.181519
(10)= 0.140259
(11)=-0.308838
(12)=-1.227539
(13)= 0.475342
(14)= 0.270996
(15)=-0.618652
(16)=-0.033936
(17)=-0.113770
(18)=-0.421387
(19)=-0.118774
(20)= 2.476562
```

Debug in Python

Debugging

- Some useful flag

- export CUTE_DSL_KEEP_IR=1
 - Keep generated CuTe IR
- export CUTE_DSL_KEEP_PTX=1
 - Keep ptx in a *.ptx file
- export CUTE_DSL_KEEP_CUBIN=1
 - Keep cubin in a *.cubin file

- Enable lineinfo to do source/sass correlation in nsight-compute

```
# Compile gemm kernel
compiled_gemm = cute.compile[cute.GenerateLineInfo(True)](
    gemm,
    a_ptr,
    b_ptr,
    sfa_ptr,
    sfb_ptr,
    c_ptr,
    (m, n, k, l),
    current_stream,
)
```

The screenshot shows two panels of the NVIDIA Nsight Compute interface. The left panel, titled 'Source', displays the CUTE-generated assembly code with line numbers and labels. The right panel, titled 'Instructions Executed', shows the corresponding GPU assembly instructions with their addresses, source lines, and assembly mnemonics. This allows for precise source-to-binary correlation during debugging.

Add new feature in Python

Inline ptx

- For cases that lack feature supports in wheel, users can inline ptx to unblock themselves

```
# /////////////////////////////////
# silu activation
# /////////////////////////////////
if self._enable_fast_sigmoid:
    t1 = acc_S.load()
    t2 = t1 * 0.5
    acc_S.store(t2)
    for i in cutlass.range_constexpr(cute.size(acc_S.shape[0])):
        for j in cutlass.range_constexpr(cute.size(acc_S.shape[1])):
            for k in cutlass.range_constexpr(cute.size(acc_S.shape[2])):
                ret = llvm.inline_asm(
                    cutlass.Float32.mlir_type,
                    [acc_S[i, j, k].ir_value()],
                    "tanh.approx.f32 $0, $1;",
                    "=f,f",
                    has_side_effects=False,
                    is_align_stack=False,
                    asm_dialect=llvm.AsmDialect.AD_ATT,
                )
                acc_S[i, j, k] = ret
    t3 = acc_S.load()
    t4 = t2 * t3 + t2
    acc_S.store(t4)
```

[hstu_attention](#)

Other Resources

- [CuTe at GPU Mode](#)
- [CuTeDSL documentations](#)
- [CuTeDSL examples](#)

THANK YOU !